# MAKING NEURAL PROGRAMMING ARCHITECTURES GENERALIZE VIA RECURSION

Jonathon Cai, Richard Shin, Dawn Song
Department of Computer Science
University of California, Berkeley
Berkeley, CA 94720, USA
{jonathon, ricshin, dawnsong}@cs.berkeley.edu

## **ABSTRACT**

Empirically, neural networks that attempt to learn programs from data have exhibited poor generalizability. Moreover, it has traditionally been difficult to reason about the behavior of these models beyond a certain level of input complexity. In order to address these issues, we propose augmenting neural architectures with a key abstraction: recursion. As an application, we implement recursion in the Neural Programmer-Interpreter framework on four tasks: grade-school addition, bubble sort, topological sort, and quicksort. We demonstrate superior generalizability and interpretability with small amounts of training data. Recursion divides the problem into smaller pieces and drastically reduces the domain of each neural network component, making it tractable to prove guarantees about the overall system's behavior. Our experience suggests that in order for neural architectures to robustly learn program semantics, it is necessary to incorporate a concept like recursion.

#### 1 Introduction

Training neural networks to synthesize robust programs from a small number of examples is a challenging task. The space of possible programs is extremely large, and composing a program that performs robustly on the infinite space of possible inputs is difficult—in part because it is impractical to obtain enough training examples to easily disambiguate amongst all possible programs. Nevertheless, we would like the model to quickly learn to represent the right semantics of the underlying program from a small number of training examples, not an exhaustive number of them.

Thus far, to evaluate the efficacy of neural models on programming tasks, the only metric that has been used is generalization of expected behavior to inputs of greater complexity (Vinyals et al. (2015), Kaiser & Sutskever (2015), Reed & de Freitas (2016), Graves et al. (2016), Zaremba et al. (2016)). For example, for the addition task, the model is trained on short inputs and then tested on its ability to sum inputs with much longer numbers of digits. Empirically, existing models suffer from a common limitation—generalization becomes poor beyond a threshold level of complexity. Errors arise due to undesirable and uninterpretable dependencies and associations the architecture learns to store in some high-dimensional hidden state. This makes it difficult to reason about what the model will do when given complex inputs.

One common strategy to improve generalization is to use curriculum learning, where the model is trained on inputs of gradually increasing complexity. However, models that make use of this strategy eventually fail after a certain level of complexity (e.g. the single-digit multiplication task in Zaremba et al. (2016), the bubble sort task in Reed & de Freitas (2016), and the graph tasks in Graves et al. (2016). In this version of curriculum learning, even though the inputs are gradually becoming more complex, the semantics of the program is succinct and does not change. Although the model is exposed to more and more data, it might learn spurious and overly complex representations of the program, as suggested in Zaremba et al. (2016). That is to say, the network does not learn the true program semantics.

In this paper, we propose to resolve these issues by explicitly incorporating recursion into neural architectures. Recursion is an important concept in programming languages and a critical tool to

reduce the complexity of programs. We find that recursion makes it easier for the network learn the right program and generalize to unknown situations. Recursion enables provable guarantees on neural programs' behavior without needing to exhaustively enumerate all possible inputs to the programs. This paper is the first (to our knowledge) to investigate the important problem of provable generalization properties of neural programs. As an application, we incorporate recursion into the Neural Programmer-Interpreter architecture and consider four sample tasks: grade-school addition, bubble sort, topological sort, and quicksort. Empirically, we observe that the learned recursive programs solve all valid inputs with 100% accuracy after training on a very small number of examples, out-performing previous generalization results. We provide proofs that these learned programs generalize perfectly.

## 2 THE PROBLEM AND OUR APPROACH

#### 2.1 The Problem of Generalization

When constructing a neural network for the purpose of learning a program, there are two orthogonal aspects to consider. The first is the actual model architecture. Numerous models have been proposed for learning programs; to name a few, this includes the Differentiable Neural Computer (Graves et al., 2016), Neural Turing Machine (Graves et al., 2014), Neural GPU (Kaiser & Sutskever, 2015), Neural Programmer (Neelakantan et al., 2015), Pointer Network (Vinyals et al., 2015), Hierarchical Attentive Memory (Andrychowicz & Kurach, 2016), and Neural Random Access Machine (Kurach et al., 2016). The architecture usually possesses some form of memory, which could be internal (such as the hidden state of a recurrent neural network) or external (such as a discrete "scratch pad" or a memory block with differentiable access). The second is the training procedure, which consists of the form of the training data and the optimization process. Almost all architectures train on program input/output pairs. The only model, to our knowledge, that does not train on input-output pairs is the Neural Programmer-Interpreter (Reed & de Freitas, 2016), which trains on synthetic execution traces.

To evaluate a neural network that learns a neural program to accomplish a certain task, one common evaluation metric is how well the learned model M generalizes. More specifically, when M is trained on simpler inputs, such as inputs of a small length, the *generalization* metric evaluates how well M will do on more complex inputs, such as inputs of much longer length. M is considered to have perfect generalization if M can give the right answer for any input, such as inputs of arbitrary length.

As mentioned in Section 1, all approaches to neural programming today fare poorly on this generalization issue. We hypothesize that the reason for this is that the neural network learns spurious dependencies which depend on specific characteristics of the training examples that are irrelevant to the true program semantics, such as length of the training inputs, and thus fails to generalize to more complex inputs.

In addition, none of the current approaches to neural programming provides a method or even aims to enable provable guarantees about generalization. The memory updates of these neural programs are so complex and interdependent that it is difficult to reason about the behaviors of the learned neural program under previously unseen situations (such as problems with longer inputs). However, this is highly undesirable, since being able to provide the correct answer in all possible settings is one of the most important aspects of any learned neural program.

## 2.2 OUR APPROACH USING RECURSION

In this paper, we propose that the key abstraction of *recursion* is necessary for neural programs to generalize. The general notion of recursion has been an important concept in many domains, including mathematics and computer science. In computer science, recursion involves solving a larger problem by combining solutions to smaller instances of the same problem (as opposed to iteration). Formally, a function exhibits recursive behavior when it possesses two properties: (1) Base cases—terminating scenarios that do not use recursion to produce answers; (2) A set of rules that reduces all other problems toward the base cases. Some functional programming languages go so far as not to define any looping constructs but rely solely on recursion to repeatedly call code.

In this paper, we propose that recursion is an important concept for neural programs as well. In fact, we argue that recursion is an essential element for neural programs to generalize, and makes it tractable to prove the generalization of neural programs. Recursion can be implemented differently for different neural programming models. Here as a concrete and general example, we consider a general Neural Programming Architecture (NPA), similar to Neural Programmer-Interpreter (NPI) in Reed & de Freitas (2016). In this architecture, we consider a core controller, e.g., an LSTM in NPI's case, but possibly other networks in different cases. There is a (changing) list of neural programs used to accomplish a given task. The core controller acts as a dispatcher for the programs. At each time step, the core controller can decide to select one of the programs to call with certain arguments. When the program is called, the current context including the caller's memory state is stored on a stack; when the program returns, the stored context is popped off the stack to resume execution in the previous caller's context.

In this general Neural Programming Architecture, we show it is easy to support recursion. In particular, recursion can be implemented as a program calling itself. Because the context of the caller is stored on a stack when it calls another program and the callee starts in a fresh context, this enables recursion simply by allowing a program to call itself. In practice, tail recursion can serve as an easy optimization when the recursive call depth is too deep. Thus, any general Neural Programming Architecture supporting such a call structure can be made to support recursion. In particular, this condition is satisfied by NPI, and thus the NPI model naturally supports recursion (even though the authors of NPI did not consider this aspect explicitly).

By nature, recursion reduces the complexity of a problem to simpler instances. Thus, recursion helps decompose a problem and makes it easier to reason about a program's behavior for previously unseen situations such as longer inputs. In particular, given that a recursion is defined by two properties as mentioned before, the base cases and the set of reduction rules, we can prove a recursive neural program generalizes perfectly if we can prove that (1) it performs correctly on the base cases; (2) it learns the reduction rules correctly. For many problems, the base cases and reduction rules usually consist of a finite (often small) number of cases. For problems where the simple cases may be infinite, such as certain forms of motor control, recursion can still help reduce the problem of generalization to these two aspects and make the generalization problem significantly simpler to handle and reason about.

As a concrete instantiation, we show in this paper that we can enable recursive neural programs in the NPI model, and thus enable perfectly generalizable neural programs for tasks such as sorting where the original NPI program fails. As aforementioned, the NPI model naturally supports recursion. However, the authors of NPI did not consider explicitly the notion of recursion and as a consequence, did not learn recursive programs. We show that by modifying the training procedure, we enable the NPI model to learn recursive neural programs. As a consequence, our learned neural programs achieve perfect generalization from a very small number of training examples. In addition, we can prove that the learned neural programs achieve perfect generalization. This is the first time one can provide provable guarantees of perfect generalization for neural programs.

We would also like to point out that in this paper, we provide as an example one way to train a recursive neural program, by providing a certain training execution trace to the NPI model. However, our concept of recursion for neural programs is general. In fact, it is one of our future directions to explore new ways to train a recursive neural program without providing explicit training execution traces or with only partial or non-recursive traces.

## 3 APPLICATION TO LEARNING RECURSIVE NEURAL PROGRAMS WITH NPI

## 3.1 BACKGROUND: NPI ARCHITECTURE

As discussed in Section 2, the Neural Programmer-Interpreter (NPI) is an instance of a Neural Programmer Architecture and hence it naturally supports recursion. In this section, we give a brief review of the NPI architecture from Reed & de Freitas (2016) as background.

We describe the details of the NPI model relevant to our contributions. We adapt machinery from the original paper slightly to fit our needs. The NPI model has three learnable components: a task-agnostic core, a program-key embedding, and domain-specific encoders that allow the NPI to operate in diverse environments.

The NPI accesses an external environment, Q, which varies according to the task. The core module of the NPI is an LSTM controller that takes as input a slice of the current external environment, via a set of pointers, and a program and arguments to execute. NPI then outputs the return probability and next program and arguments to execute. Formally, the NPI is represented by the following set of equations:

$$s_{t} = f_{enc}(e_{t}, a_{t})$$

$$h_{t} = f_{lstm}(s_{t}, p_{t}, h_{t-1})$$

$$r_{t} = f_{end}(h_{t}), p_{t+1} = f_{prog}(h_{t}), a_{t+1} = f_{arg}(h_{t})$$

t is a subscript denoting the time-step;  $f_{enc}$  is a domain-specific encoder (to be described later) that takes in the environment slice  $e_t$  and arguments  $a_t$ ;  $f_{lstm}$  represents the core module, which takes in the state  $s_t$  generated by  $f_{enc}$ , a program embedding  $p_t \in \mathbb{R}^P$ , and hidden LSTM state  $h_t$ ;  $f_{end}$  decodes the return probability  $r_t$ ;  $f_{prog}$  decodes a program key embedding  $p_{t+1}$ ; and  $f_{arg}$  decodes arguments  $a_{t+1}$ . The outputs  $r_t$ ,  $p_{t+1}$ ,  $a_{t+1}$  are used to determine the next action, as described in Algorithm 1. If the program is primitive, the next environmental state  $e_{t+1}$  will be affected by  $p_t$  and  $a_t$ , i.e.  $e_{t+1} \sim f_{env}(e_t, p_t, a_t)$ . As with the original NPI architecture, the experiments for this paper always used a 3-tuple of integers  $a_t = (a_t(1), a_t(2), a_t(3))$ .

## **Algorithm 1** Neural programming inference

```
1: Inputs: Environment observation e, program p, arguments a, stop threshold \alpha
 2: function RUN(e, p, a)
          h \leftarrow \mathbf{0}, r \leftarrow 0
 3:
 4:
          while r < \alpha do
 5:
               s \leftarrow f_{enc}(e, a), h \leftarrow f_{lstm}(s, p, h)
 6:
               r \leftarrow f_{end}(h), p_2 \leftarrow f_{prog}(h), a_2 \leftarrow f_{arg}(h)
               if p is a primitive function then
 7:
 8:
                    e \leftarrow f_{env}(e, p, a).
 9:
               else
10:
                    function RUN(e, p_2, a_2)
```

A description of the inference procedure is given in Algorithm 1. Each step during an execution of the program does one of three things: (1) another subprogram along with associated arguments is called, as in Line 10, (2) the program writes to the environment if it is primitive, as in Line 8, or (3) the loop is terminated if the return probability exceeds a threshold  $\alpha$ , after which the stack frame is popped and control is returned to the caller. In all experiments,  $\alpha$  is set to 0.5. Each time a subprogram is called, the stack depth increases.

The training data for the Neural Programmer-Interpreter consists of full execution traces for the program of interest. A single element of an execution trace consists of a step input-step output pair, which can be synthesized from Algorithm 1: this corresponds to, for a given time-step, the step input tuple (e,p,a) and step output tuple  $(r,p_2,a_2)$ . An example of part of an addition task trace, written in shorthand, is given in Figure 1. For example, a step input-step output pair in Lines 2 and 3 of the left-hand side of Figure 1 is (ADD1, WRITE OUT 1). In this pair, the step input runs a subprogram ADD1 that has no arguments, and the step output contains a program WRITE that has arguments of OUT and 1. The environment and return probability are omitted for readability. Indentation indicates the stack is one level deeper than before.

It is important to emphasize that at inference time in the NPI, the hidden state of the LSTM controller is reset (to zero) at each subprogram call, as in Line 3 of Algorithm 1 ( $h \leftarrow \mathbf{0}$ ). This functionality is critical for implementing recursion, since it permits us to restrict our attention to the currently relevant recursive call, ignoring irrelevant details about other contexts.

<sup>&</sup>lt;sup>1</sup>The original NPI paper decodes to a program key embedding  $k_t \in \mathbb{R}^K$  and then computes a program embedding  $p_{t+1}$ , which we also did in our implementation, but we omit this for brevity.

#### 3.2 RECURSIVE FORMULATIONS FOR NPI PROGRAMS

We emphasize the overall goal of this work is to enable the learning of a recursive program. The learned recursive program is different from neural programs learned in all previous work in an important aspect: previous approaches do not explicitly incorporate this abstraction, and hence generalize poorly, whereas our learned neural programs incorporate recursion and achieve perfect generalization.

Since NPI naturally supports the notion of recursion, a key question is how to enable NPI to learn recursive programs. We found that changing the NPI training traces is a simple way to enable this. In particular, we construct new training traces which explicitly contain recursive elements and show that with this type of trace, NPI easily learns recursive programs. In future work, we would like to decrease supervision and construct models that are capable of coming up with recursive abstractions themselves.

In what follows, we describe the way in which we constructed NPI traces so as to make them contain recursive elements and thus enable NPI to learn recursive programs. We describe the recursive reformulation of traces for two tasks from the original NPI paper—grade-school addition and bubble sort. For these programs, we re-use the appropriate program sets (the associated subprograms), and we refer the reader to the appendix of Reed & de Freitas (2016) for further details on the subprograms used in addition and bubble sort. Finally, we implement recursive traces for our own topological sort and quicksort tasks.

Grade School Addition. For grade-school addition, the domain-specific encoder is

$$f_{enc}(Q, i_1, i_2, i_3, i_4, a_t) = MLP([Q(1, i_1), Q(2, i_2), Q(3, i_3), Q(4, i_4), a_t(1), a_t(2), a_t(3)]),$$

where the environment  $Q \in \mathbb{R}^{4 \times N \times K}$  is a scratch-pad that contains four rows (the first input number, the second input number, the carry bits, and the output) and N columns. K is set to 11, to represent the range of 10 possible digits, along with the end token.<sup>2</sup> At any given time, the NPI has access to four pointers in each of the four rows, represented by  $Q(1,i_1), Q(2,i_2), Q(3,i_3)$ , and  $Q(4,i_4)$ .

The non-recursive trace loops on cycles of ADD1 and LSHIFT. ADD1 is a subprogram that adds the current column (writing the appropriate digit to the output row and carrying a bit to the next column if needed). LSHIFT moves the four pointers to the left, to move to the next column. The program terminates when seeing no numbers in the current column.

Figure 1 shows examples of non-recursive and recursive addition traces. We make the trace recursive by adding a tail recursive call into the trace for the ADD program after calling ADD1 and LSHIFT, as in Line 13 of the right-hand side of Figure 1. Via the recursive call, we effectively forget that the column just added exists, since we reset the hidden state of the LSTM controller upon the recursive call to ADD. Consequently, there is no concept of length relevant to the problem, which is traditionally important (as in length-based curriculum learning).

**Bubble Sort.** For bubble sort, the domain-specific encoder is

$$f_{enc}(Q, i_1, i_2, i_3, a_t) = MLP([Q(1, i_1), Q(1, i_2), Q(1, i_3), a_t(1), a_t(2), a_t(3)]),$$

where the environment  $Q \in \mathbb{R}^{1 \times N \times K}$  is a scratch-pad that contains 1 row, to represent the state of the array as sorting proceeds in-place, and N columns. K is set to 11, to denote the range of possible numbers (0 through 9), along with the start/end token (represented with the same encoding). At any given time, the NPI has access to three pointers, represented by  $Q(1,i_1),Q(2,i_2)$ , and  $Q(3,i_3)$ . The pointers at index  $i_1$  and  $i_2$  are used to compare the pair of numbers considered during the bubble sweep, swapping them if the number at  $i_1$  is greater than that in  $i_2$ . These pointers are referred to as bubble pointers. The pointer at index  $i_3$  represents a counter internal to the environment that is incremented once after each pass of the algorithm (one cycle of BUBBLE and RESET); it terminates the entire algorithm when incremented a number of times equal to the length of the array. We restrict  $Q(3,i_3)$  to  $\{0,1\}$ , since it is simply a flag that checks whether or not the end of the array has been reached.

<sup>&</sup>lt;sup>2</sup>The original paper uses K=10, but we found it necessary to augment the range with an end token, in order to terminate properly.

#### Non-Recursive Recursive ADD1 ADD1 WRITE OUT 1 WRITE OUT 1 CARRY CARRY PTR CARRY LEFT PTR CARRY LEFT WRITE CARRY 1 WRITE CARRY 1 PTR CARRY RIGHT PTR CARRY RIGHT LSHIFT LSHIFT PTR INP1 LEFT PTR INP1 LEFT 10 PTR INP2 LEFT 10 PTR INP2 LEFT PTR CARRY LEFT PTR CARRY LEFT PTR OUT LEFT PTR OUT LEFT 13 ADD1 13 ADD

Figure 1: Addition Task. The non-recursive trace loops on cycles of ADD1 and LSHIFT, whereas in the recursive version, the ADD function calls itself (bolded).

	Non-Recursive	Partial Recursive	Full Recur	sive
1 2 3 4 4 5 6 6 7 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23	BUBBLESORT BUBBLE PTR 2 RIGHT BSTEP COMPSWAP  RSHIFT PTR 1 RIGHT PTR 2 RIGHT BSTEP COMPSWAP SWAP 1 2 RSHIFT PTR 1 RIGHT PTR 2 RIGHT PTR 2 RIGHT LSHIFT PTR 1 LEFT PTR 1 LEFT PTR 2 LEFT LSHIFT PTR 1 LEFT PTR 2 LEFT LSHIFT PTR 1 LEFT PTR 2 LEFT PTR 2 LEFT PTR 3 RIGHT	1 BUBBLESORT 2 BUBBLE 3 PTR 2 RIGHT 4 BSTEP 5 COMPSWAP 6 7 RSHIFT 8 PTR 1 RIGHT 9 PTR 2 RIGHT 10 BSTEP 11 COMPSWAP 12 SWAP 1 2 13 RSHIFT 14 PTR 1 RIGHT 15 PTR 2 RIGHT 16 RESET 17 LSHIFT 18 PTR 1 RIGHT 16 RESET 17 LSHIFT 18 PTR 1 LEFT 19 PTR 2 LEFT 20 LSHIFT 21 PTR 1 LEFT 22 PTR 2 LEFT 23 PTR 3 RIGHT 24 BUBBLESORT	1 BUBBLESORT 2 BUBBLE 3 PTR 2 R1 4 BSTEP 5 COMPSV 6 7 RSHIFT 8 PTR 9 PTR 10 BSTEP 11 COMM 12 SW 13 RSH1 14 PT 15 PT 16 RESET 17 LSHIFT 18 PTR 1 19 PTR 2 20 LSHIFT 21 PTR	IGHT  NAP  1 RIGHT 2 RIGHT 2 RIGHT PSWAP 1 2 IFT RR 1 RIGHT RR 2 RIGHT LEFT LEFT 1 LEFT 2 LEFT 1 LEFT 1 LEFT
24 25	BUBBLE	25 BUBBLE 26	25 BUBBLE 26	

Figure 2: Bubble Sort Task. The non-recursive trace loops on cycles of BUBBLE and RESET. The difference between the partial recursive and full recursive versions is in the indentation of Lines 10-15 and 20-22 (bolded), since in the full recursive version, BSTEP and LSHIFT are made tail recursive. Also note that COMPSWAP conditionally swaps numbers under the bubble pointers.

The non-recursive trace loops on cycles of BUBBLE and RESET, which logically represents one bubble sweep through the array and reset of the two bubble pointers to the very beginning of the array, respectively. In this version, there is a dependence on length: BSTEP and LSHIFT are called a number of times equivalent to one less than the length of the input array, in BUBBLE and RESET respectively.

Inside BUBBLE and RESET, there are two operations that can be made recursive. BSTEP, used in BUBBLE, compares pairs of numbers, continuously moving the bubble pointers once to the right each time until reaching the end of the array. LSHIFT, used in RESET, shifts the pointers left until reaching the start token.

We experiment with two levels of recursion—partial and full. Partial recursion only adds a tail recursive call to BUBBLESORT after BUBBLE and RESET, similar to the tail recursive call described previously for addition. The partial recursion is not enough for perfect generalization, as will be presented later in Section 4. Full recursion, in addition to making the aforementioned tail recursive call, adds two additional recursive calls; BSTEP and LSHIFT are made tail recursive. Figure 2 shows examples of traces for the different versions of bubble sort. Training on the full recursive trace leads to perfect generalization, as shown in Section 4. We performed experiments on the partially recursive version in order to examine what happens when only one recursive call is implemented, when in reality three are required for perfect generalization.

## Algorithm 2 Depth First Search Topological Sort

```
1: Color all vertices white.
 2: Initialize an empty stack S and a directed acyclic graph DAG to traverse.
 3: Begin traversing from Vertex 1 in the DAG.
 4: function TopoSort(DAG)
 5:
        while there is still a white vertex u: do
 6:
            color[u] = grey
 7:
            v_{active} = u
 8:
            do
 9:
                if v_{active} has a white child v then
10:
                    color[v] = grey
                    push v_{active} onto S
11:
                    v_{active} = v
12:
                else
13:
14:
                    color[v_{active}] = black
15:
                    Write v_{active} to result
                    if S is empty then pass
16:
                    else pop the top vertex off S and set it to v_{active}
17:
18:
            while S is not empty
```

**Topological Sort.** We choose to implement a topological sort task for graphs. A *topological sort* is a linear ordering of vertices such that for every directed edge (u, v) from u to v, u comes before v in the ordering. This is possible if and only if the graph has no directed cycles; that is to say, it must be a directed acyclic graph (DAG). In our experiments, we only present DAG's as inputs and represent the vertices as values ranging from  $1, \ldots, n$ , where the DAG contains n vertices.

Directed acyclic graphs are structurally more diverse than inputs in the two tasks of grade-school addition and bubble sort. The degree for any vertex in the DAG is variable. Also the DAG can have potentially more than one connected component, meaning it is necessary to transition between these components appropriately.

Algorithm 2 shows the topological sort task of interest. This algorithm is a variant of depth first search. We created a program set that reflects the semantics of Algorithm 2. For brevity, we refer the reader to the appendix for further details on the program set and non-recursive and recursive trace-generating functions used for topological sort.

For topological sort, the domain-specific encoder is

```
f_{enc}(DAG, Q_{color}, p_{stack}, p_{start}, v_{active}, childList)
= MLP([Q_{color}(p_{start}), Q_{color}(DAG[v_{active}][childList[v_{active}]]), p_{stack} == 1, a_t(1), a_t(2), a_t(3)]),
```

where  $Q_{color} \in \mathbb{R}^{U \times 4}$  is a scratch-pad that contains U rows, each containing one of four colors (white, gray, black, invalid) with one-hot encoding. U varies with the number of vertices in the graph. We further have  $Q_{result} \in \mathbb{N}^U$ , a scratch-pad which contains the sorted list of vertices at the end of execution, and  $Q_{stack} \in \mathbb{N}^U$ , which serves the role of the stack S in Algorithm 2. The contents of  $Q_{result}$  and  $Q_{stack}$  are not exposed directly through the domain-specific encoder; rather, we define primitive functions which manipulate these scratch-pads.

The DAG is represented as an adjacency list where DAG[i][j] refers to the j-th child of vertex i. There are 3 pointers  $(p_{result}, p_{stack}, p_{start})$ ,  $p_{result}$  points to the next empty location in  $Q_{result}$ ,  $p_{stack}$  points to the top of the stack in  $Q_{stack}$ , and  $p_{start}$  points to the candidate starting node for a connected component. There are 2 variables  $(v_{active} \text{ and } v_{save})$ ;  $v_{active}$  holds the active vertex (as in Algorithm 2) and  $v_{save}$  holds the value of  $v_{active}$  before executing Line 12 of Algorithm 2.  $childList \in \mathbb{N}^U$  is a vector of pointers, where childList[i] points to the next child under consideration for vertex i.

The three environment observations aid with control flow in Algorithm 2.  $Q_{color}(p_{start})$  contains the color of the current start vertex, used in the evaluation of the condition in the WHILE loop in Line 5 of Algorithm 2.  $Q_{color}(DAG[v_{active}][childList[v_{active}]])$  refers to the color of the next child of

 $v_{active}$ , used in the evaluation of the condition in the IF branch in Line 9 of Algorithm 2. Finally, the boolean  $p_{stack} == 1$  is used to check whether the stack is empty in Line 18 of Algorithm 2.

An alternative way of representing the environment slice is to expose the values of the absolute vertices to the model; however, this makes it difficult to scale the model to larger graphs, since large vertex values are not seen during training time.

We refer the reader to the appendix for the non-recursive trace generating functions. In the non-recursive trace, there are four functions that can be made recursive—TOPOSORT, CHECK\_CHILD, EXPLORE, and NEXT\_START, and we add a tail recursive call to each of these functions in order to make the recursive trace. In particular, in the EXPLORE function, adding a tail recursive call resets and stores the hidden states associated with vertices in a stack-like fashion. This makes it so that we only need to consider the vertices in the subgraph that are currently relevant for computing the sort, allowing simpler reasoning about behavior for large graphs. The sequence of primitive operations (MOVE and WRITE operations) for the non-recursive and recursive versions are exactly the same.

**Quicksort.** We implement a quicksort task, in order to demonstrate that recursion helps with learning divide-and-conquer algorithms. We use the Lomuto partition scheme; the logic for the recursive trace is shown in Algorithm 3. For brevity, we refer the reader to the appendix for information about the program set and non-recursive and recursive trace-generating functions for quicksort. The logic for the non-recursive trace is shown in Algorithm 4 in the appendix.

## **Algorithm 3** Recursive Quicksort

```
1: Initialize an array A to sort.
2: Initialize lo and hi to be 1 and n, where n is the length of A.
4: function QUICKSORT(A, lo, hi)
5:
        if lo < hi: then
6:
            p = PARTITION(A, lo, hi)
7:
            QUICKSORT(A, lo, p - 1)
8:
            QUICKSORT(A, p + 1, hi)
9:
10: function PARTITION(A, lo, hi)
        pivot = lo
11:
        \textbf{for } j \in [lo, hi-1]: \textbf{do}
12:
13:
            if A[j] \leq A[hi] then
                swap A[pivot] with A[j]
14:
15:
               pivot = pivot + 1
        swap A[pivot] with A[hi]
16:
17:
        return pivot
```

For quicksort, the domain-specific encoder is

```
\begin{split} f_{enc}(Q_{array},Q_{stackLo},Q_{stackHi},p_{lo},p_{hi},p_{stackLo},p_{stackHi},p_{pivot},p_{j},a_{t}) = \\ MLP([Q_{array}(p_{j}) < Q_{array}(p_{hi}),p_{j} == p_{hi},p_{pivot} == p_{lo},p_{j} == p_{lo},p_{pivot} == 1,p_{j} == 1,\\ Q_{stackLo}(p_{stackLo} - 1) < Q_{stackHi}(p_{stackHi} - 1),p_{stackLo} == 1,a_{t}(1),a_{t}(2),a_{t}(3)]), \end{split}
```

where  $Q_{array} \in \mathbb{R}^{U \times 11}$  is a scratch-pad that contains U rows, each containing one of 11 values (one of the numbers 0 through 9 or an invalid state). Our implementation uses two stacks  $Q_{stackLo}$  and  $Q_{stackHi}$ , each in  $\mathbb{R}^U$ , that store the arguments to the recursive QUICKSORT calls in Algorithm 3; before each recursive call, the appropriate arguments are popped off the stack and written to  $p_{lo}$  and  $p_{hi}$ .

There are 6 pointers  $(p_{lo}, p_{hi}, p_{stackLo}, p_{stackHi}, p_{pivot}, p_j)$ .  $p_{lo}$  and  $p_{hi}$  point to the lo and hi indices of the array, as in Algorithm 3.  $p_{stackLo}$  and  $p_{stackHi}$  point to the top (empty) positions in  $Q_{stackLo}$  and  $Q_{stackHi}$ .  $p_{pivot}$  and  $p_j$  point to the pivot and j indices of the array, used in the PARTITION function in Algorithm 3. The 8 environment observations aid with control flow;  $Q_{stackLo}(p_{stackLo}-1) < Q_{stackHi}(p_{stackHi}-1)$  implements the lo < hi comparison in Line 5 of

Algorithm 3,  $p_{stackLo} == 1$  checks if the stacks are empty in Line 18 of Algorithm 4, and the other observations (all involving  $p_{pivot}$  or  $p_i$ ) deal with logic in the PARTITION function.

Note that the recursion for quicksort is not purely tail recursive and therefore represents a more complex kind of recursion that is harder to learn than in the previous tasks. Also, compared to the bubble pointers in bubble sort, the pointers that perform the comparison for quicksort (the COMP-SWAP function) are usually not adjacent to each other, making quicksort less local than bubble sort. In order to compensate for this,  $p_{pivot}$  and  $p_j$  require special functions (MOVE\_PIVOT\_LO, MOVE\_J\_LO, RESET\_PIVOT, and RESET\_J) to properly set them to lo in Lines 11 and 12 of the PARTITION function in Algorithm 3. For the non-recursive version of the program, these functions can introduce spurious dependencies on length (specific to the training inputs) into the model.

#### 3.3 Provably Perfect Generalization

We show that if we incorporate recursion, the learned NPI programs achieve provably perfect generalization. Provably perfect generalization implies the model will behave correctly, given any valid input. In order to claim a proof, we must verify the model produces correct behavior over all base cases and reductions, as described in Section 2.

We describe our verification procedure. This procedure verifies that all base cases and reductions are handled properly by the model via explicit tests. Note that recursion helps make this process tractable, because it reduces the number of inputs that need to be tested. This verification phase only needs to be performed once after training.

Formally, verification consists of proving the following theorem:

$$\forall i \in V(S), M(i) \downarrow P(i)$$

where i denotes a step input, V(S) denotes the set of valid step inputs, M denotes the neural network model, P denotes the correct program, and P(i) denotes the output from the program. The arrow in the theorem refers to evaluation, as in big-step semantics. The theorem states that for the same input, the model produces the exact same output as the target program it aims to learn.

Recursion drastically reduces the number of configurations we need to consider during the verification phase and makes the proof tractable, because it introduces structure that eliminates infinite families of inputs that would otherwise need to be considered. For instance, for recursive bubble sort, we prove the family F of arrays consisting of only 2's and 3's (of any length) is sorted properly given that all members of  $S = \{[2,2],[2,3],[3,2],[3,3]\}$  are sorted correctly. Without recursion, such a proof is not possible, because we would not be able to easily reason about the behavior of problems in F in terms of subproblems in S. In the non-recursive formulation, BUBBLE calls BSTEP a number of times that is dependent on the length of the input. The core LSTM module's hidden state is preserved over all these BSTEP calls, making it difficult to interpret with certainty what happens over longer timesteps. In contrast, if the context is refreshed as in the recursive BSTEP, when sorting any problem in F, the states from the BSTEP calls from the 4 subproblems in S can be reused. This guarantees that subproblem outputs generated during verification are generated during execution of an unseen problem in F, leading to generalization to any problem in F. Hence, if all subproblems in S are sorted correctly, we have proven that any member of F will be sorted correctly, thus eliminating an infinite family of inputs that need to be tested.

Now that we have described the verification procedure, we describe the space of base cases and reductions that must be covered for each of the four sample tasks, in order to claim a proof of correctness.

**Base Cases and Reduction for Addition.** For the recursive formulation of addition, the base cases are the sequences associated with one cycle of ADD, which means ensuring that (1): ADD1 computes the current column's sum properly (and carries, if necessary), and (2): LSHIFT moves all four pointers successfully to the next column. It must be proven that ADD1 sums every possible triple of single digits in a column (including the carry bit) properly: there are 2\*10\*10=200 such sums, excluding the end state. Moreover, in the process of moving to the next column, a *partial state* arises that spans two columns: since each pointer is moved individually, there can arise partial states where at least one pointer is in the current column and at least one other pointer is in the next

column. For each of the four pointer movements of LSHIFT, the number of partial states to consider is  $2*10^3$ , excluding the end state. In order to prove that LSHIFT works properly, it is necessary to prove that under all these partial states, LSHIFT behaves as expected. Observe that the network should learn to LSHIFT by ignoring these partial states, since LSHIFT moves each pointer to the left regardless of the partial state values. Nevertheless, when proving that the algorithm works, we must test over all these partial states. This is a limitation of the current NPI model, and in future work we hope to consider how to reduce the number of partial states that need to be tested during the verification procedure. Finally, the algorithm must terminate when not seeing any numbers in a column, which is so when every element is an end token. There is only one reduction rule, since there is only one tail recursive call to ADD.

**Base Cases and Reductions for Bubble Sort.** For the full recursive formulation of bubble sort, the base cases are the sequences associated with one cycle of BUBBLESORT, which means ensuring that (1): BUBBLE performs a sweep across the entire array and swaps, if necessary, the pair of numbers considered under the bubble pointers, and (2): RESET moves the bubble pointers back to the start of the array and moves the counter pointer once to the right. It must be proven that BSTEP demonstrates proper swapping behavior under every possible environment, of which there are 10\*11\*2 possible. It is especially important to prove that BSTEP terminates properly, namely when the second bubble pointer has reached an end token. A similar analysis can be applied to LSHIFT. Also, one must prove that the first time BUBBLESORT encounters an environment where the counter pointer has reached an end token, the entire algorithm terminates. There are three reduction rules, since there are three tail recursive calls to BUBBLESORT, BSTEP, and LSHIFT.

Base Cases and Reductions for Topological Sort. For the recursive formulation of topological sort, the base cases are the sequences associated with one cycle of TOPOSORT, including reaching the start of another connected component. We must account for all colors that might appear in the environment during the execution trace. For example, during the traversal to the starting node of the next connected component (during execution of NEXT\_START), it is necessary to account for possibly white, black, and invalid-colored candidate starting nodes. Additionally, during exploration of the graph, it is necessary to account for possibly white, black, grey, and invalid-colored children of  $v_{active}$ . There are four reduction rules, since there are four tail recursive calls to TOPOSORT, CHECK\_CHILD, EXPLORE, and NEXT\_START.

Base Cases and Reductions for Quicksort. For the recursive formulation of quicksort, the base cases are the sequences associated with one cycle of QUICKSORT, including proper behavior for the PARTITION function and the recursive QUICKSORT calls. Most of the environment observations are associated with only one program in the program set. For example, the observation  $p_{pivot} == 1$  is associated with the RESET\_PIVOT call. But it is still necessary to consider all possible valid combinations of the other observations, such as  $p_{pivot} == p_{lo}$ —this environment observation might be set to true when resetting  $p_{pivot}$ . Ideally, the learned program learns to ignore observations that are irrelevant to the currently executed program. There are two reduction rules, since there are two recursive calls to QUICKSORT.

**Verification over an Appropriate Set of Inputs** One could list all possible valid steps and verify them individually by feeding each singular step input into the model and affirming correctness; however, this is tedious. It is enough to show that the model exhibits correctness over an appropriate test set of problem inputs that covers the entire space of base cases and reductions.

For the addition task, it is enough to test over all 2-digit problems, of which there are 8100. Together, these test cases cover all base cases (including the unique sums) and reduction rules.

For bubble sort, it is enough to test over all possible arrays containing 2 digits, of which there are 100. Together, these test cases cover all base cases (including all possible pairs of numbers considered under the swap operation) and reduction rules.

For topological sort, it is enough to test that a particular graph of size 6 (described in the appendix) is handled properly. The test set must contain graphs that handle all possible color configurations. Only including certain graph structures in the test set, such as linear chains, is not sufficient for observing all possible configurations—for linear chains, for example, there are no black candidate starting nodes for connected components.

For quicksort, it is enough to test that a particular set of 3 arrays (described in the appendix) is sorted properly.

One can confirm that all possible valid environment and argument combinations, as well as program transitions, are present in these test sets; these test sets cover all base cases and reduction rules. Thus, a model (trained on recursive traces) that passes these tests generalizes perfectly to all valid inputs.

We call a test set a *verification set* if it admits a proof of correctness for a particular task. The verification sets described previously for each task were used in Section 4.

We note that for tasks with very large input domains, such as ones involving MNIST digits or speech samples, the state space of base cases and reduction rules could be prohibitively large, possibly infinite. Consequently, it is infeasible to construct a verification set that covers all cases, and the verification procedure we have described is inadequate. In future work, we hope to devise a verification procedure more appropriate to this setting.

#### 4 EXPERIMENTS

As there is no public implementation of NPI, we implemented a version of it in Keras that is as faithful to the paper as possible. Our experiments use a very small number of training examples. The training set for the addition task contains 200 traces in total, accounting for each unique sum mentioned in Section 3.3. The maximum problem length in this training set is 3 (e.g., the trace corresponding to the addition problem "109 + 101"). Our training set for bubble sort contains 100 traces, with a maximum problem length of 2 (e.g., the trace corresponding to the sort of the array [3,2]). The training set for bubble sort matches the verification set described in Section 3.3. Our training set for topological sort contains 1 trace, synthesized from a graph with 5 vertices (described in the appendix). The training set for quicksort matches the verification set described in Section 3.3. The same set of problems was used to generate the training traces for all formulations of the task, for non-recursive and recursive versions.

We train using the Adam optimizer and use a 2-layer LSTM and task-specific state encoders for the external environments, as described in Reed & de Freitas (2016).

We now report on generalization for the varying tasks.

#### 4.1 Grade-School Addition

Both the non-recursive and recursive learned programs generalize on all input lengths we tried, up to 5000 digits. This agrees with the generalization of non-recursive addition in Reed & de Freitas (2016), where they reported generalization up to 3000 digits, but with a larger amount of training data (640 examples, 32 each of length 1 through 20).

In order to demonstrate a setting where recursion improves the task, for addition, we trained only on traces for 5 arbitrarily chosen 1-digit addition sum examples. The recursive version can generalize perfectly to long problems constructed from these components (such as the sum "822+233", where "8+2" and "2+3" are in the training set), but the non-recursive version fails to sum these long problems properly.

#### 4.2 Bubble Sort

Table 1 presents results on randomly generated arrays of varying length for the learned non-recursive, partially recursive, and full recursive programs. For each length, we test each program on 30 randomly generated problems. Observe that partially recursive does slightly better than non-recursive for the setting in which the length of the array is 3, and that the fully recursive version is able to sort every array given to it. The non-recursive and partially recursive versions are unable to sort long arrays, beyond length 8.

Table 1: Accuracy on Randomly Generated Problems for Bubble Sort

Length of Array	Non-Recursive	<b>Partially Recursive</b>	<b>Full Recursive</b>
_			
2	100%	100%	100%
3	6.7%	23%	100%
4	10%	10%	100%
8	0%	0%	100%
20	0%	0%	100%
90	0%	0%	100%

Table 2: Accuracy on Randomly Generated Problems for Topological Sort

Number of Vertices	Non-Recursive	Recursive
-	(70	1000/
5	6.7%	100%
6	6.7%	100%
7	3.3%	100%
8	0%	100%
70	0%	100%

#### 4.3 TOPOLOGICAL SORT

Table 2 presents results on randomly generated DAG's of varying graph sizes (varying in the number of vertices) after training on just a single execution trace for the topological sort task. For each graph size, we test the learned programs on 30 randomly generated DAG's. The recursive version of topological sort solves all graph instances we tried, from graphs of size 5 through 70. On the other hand, the non-recursive version has low accuracy, beginning from size 5, and fails completely for graphs of size 8 and beyond.

## 4.4 QUICKSORT

Table 3 presents results on randomly generated arrays of varying length for the learned non-recursive and recursive programs. For each length, we test each program on 30 randomly generated problems. Observe that the non-recursive program cannot sort arrays of length 4 and arrays of length 30 and beyond, while the recursive program can sort any given array.

As mentioned in Section 2.1, we hypothesize the non-recursive programs do not generalize well because they have learned spurious dependencies specific to the training set, such as length of the input problems. On the other hand, the recursive programs have learned the true program semantics.

## 4.5 Provably Perfect Generalization

We prove that the models trained with recursive traces generalize, via the verification procedure described in Section 3.3. As described in the verification procedure, it is possible to prove our learned program, trained from recursive traces, generalizes perfectly by testing on an appropriate set of problem inputs. Recall that this verification procedure cannot be performed for the non-recursive versions, since the propagation of the hidden state in the core LSTM module makes reasoning difficult.

**Proof of Addition.** Although the authors of Reed & de Freitas (2016) state their version of non-recursive addition works for problem sizes of up to 3000, they do not have an explicit proof that their learned program is truly correct. In contrast, we are able to prove our learned program generalizes perfectly. As described in Section 3.3, for the grade-school addition task, it is possible to prove correctness of our learned program by testing on all 8100 2-digit problem inputs ("10+10", ..., "99+99"); this verification set covers the entire set of base cases and reduction rules. We explicitly test our model on all 2-digit problems and prove the learned program generalizes perfectly by ob-

Table 3: Accuracy on Randomly Generated Problems for Quicksort

Length of Array	Non-Recursive	Recursive
4	0%	100%
5	13.3%	100%
6	20%	100%
7	36.7%	100%
8	20%	100%
9	23.3%	100%
10	20%	100%
11	23.3%	100%
30	0%	100%
70	0%	100%

serving that correct traces are produced for all these inputs. It is worth noting that not all base cases are seen in the training set. We train on a set of 200 traces, insufficient to cover all possible partial states (described in Section 3.3). Nevertheless, the trained model has learned the correct program semantics; intuitively, the learned program has learned to ignore partial states when performing LSHIFT.

**Proof of Bubble Sort.** We also report substantial improvements on generalization of the bubble sort task in Reed & de Freitas (2016). In Reed & de Freitas (2016), after training the Neural Programmer-Interpreter on 1216 traces (64 each for lengths 2 through 20), generalization fails for arrays with 60 numbers and beyond. In contrast, we use only 100 traces, training on arrays with 2 numbers, and achieve provably perfect generalization on all observed problem inputs. As described in Section 3.3, for the bubble sort task, it is possible to prove correctness of our learned program by testing on all 100 arrays with 2 numbers ([0,0], [0,1], ..., [9,9]); this verification set covers the entire set of base cases and reduction rules. We explicitly test our model on all arrays with 2 numbers and prove the learned program generalizes perfectly by observing that correct traces are produced for all these inputs.

**Proof of Topological Sort.** We prove perfect generalization for the model trained on a single recursive trace for the topological sort task. As described in Section 3.3, for the topological sort task, it is possible to prove correctness of our learned program by verifying correctness on a particular graph of size 6 (described in the appendix); this problem input covers the entire set of base cases and reduction rules. We explicitly test our model on the aforementioned graph and prove the learned program generalizes perfectly by observing the correct trace is produced for this input. It is worth noting that not all base cases are seen in the training trace, in particular the setting of a single vertex in a connected component (a *lone vertex*). We train our model on a graph of size 5 composed of two connected components, where each component has at least two vertices and hence no lone vertices. Nevertheless, the learned program handles the setting of lone vertices correctly. The trained model has learned the correct program semantics; intuitively, in the case of lone vertices, the model has learned to execute the correct program sequence before calling the TRAVERSE method (this sequence does not rely on the second environment observation—color of child of  $v_{active}$ —exposed to the model).

**Proof of Quicksort.** We prove perfect generalization for the model trained on recursive traces for the quicksort task. As described in Section 3.3, for the quicksort task, it is possible to prove correctness of our learned program by verifying correctness on a set of size 3 (described in the appendix); these input arrays cover the entire set of base cases and reduction rules. We explicitly test our model on these arrays and prove the learned program generalizes perfectly by observing that correct traces are produced for all these inputs.

We demonstrate that recursion enables proving perfect generalization for different tasks, including addition, bubble sort, topological sort, and quicksort. Empirically, we observe perfect generalization on all inputs given to models trained with recursive traces.

### 5 CONCLUSION

We emphasize that the notion of a neural recursive program has not been presented in the literature before: this is our main contribution. Recursion enables provably perfect generalization. To our knowledge, this is the first time verification has been applied to a neural network, providing provable guarantees about its behavior. We instantiated recursion for the Neural Programmer-Interpreter by changing the training traces. In future work, we seek to implement more tasks with recursive structure. We also hope to decrease supervision, perhaps by training with only partial or non-recursive traces, and to integrate a notion of recursion into the models themselves by constructing novel Neural Programming Architectures.

#### ACKNOWLEDGMENTS

This material is in part based upon work supported by the National Science Foundation under Grant No. TWC-1409915 and DARPA under Grant No. FA8750-15-2-0104. Any opinions, findings, and conclusions or recommendations expressed in this material are those of the author(s) and do not necessarily reflect the views of National Science Foundation and DARPA.

## REFERENCES

- Marcin Andrychowicz and Karol Kurach. Learning efficient algorithms with hierarchical attentive memory. *CoRR*, abs/1602.03218, 2016. URL http://arxiv.org/abs/1602.03218.
- Alex Graves, Greg Wayne, and Ivo Danihelka. Neural turing machines. *CoRR*, abs/1410.5401, 2014. URL http://arxiv.org/abs/1410.5401.
- Alex Graves, Greg Wayne, Malcolm Reynolds, Tim Harley, Ivo Danihelka, Agnieszka Grabska-Barwiska, Sergio Gmez Colmenarejo, Edward Grefenstette, Tiago Ramalho, John Agapiou, Adri Puigdomnech Badia, Karl Moritz Hermann, Yori Zwols, Georg Ostrovski, Adam Cain, Helen King, Christopher Summerfield, Phil Blunsom, Koray Kavukcuoglu, and Demis Hassabis. Hybrid computing using a neural network with dynamic external memory. *Nature*, 538 (7626):471–476, October 2016. ISSN 0028-0836, 1476-4687. doi: 10.1038/nature20101. URL http://www.nature.com/doifinder/10.1038/nature20101.
- Lukasz Kaiser and Ilya Sutskever. Neural gpus learn algorithms. *CoRR*, abs/1511.08228, 2015. URL http://arxiv.org/abs/1511.08228.
- Karol Kurach, Marcin Andrychowicz, and Ilya Sutskever. Neural random access machines. *ERCIM News*, 2016(107), 2016. URL http://ercim-news.ercim.eu/en107/special/neural-random-access-machines.
- Arvind Neelakantan, Quoc V. Le, and Ilya Sutskever. Neural programmer: Inducing latent programs with gradient descent, 2015.
- Scott Reed and Nando de Freitas. Neural programmer-interpreters. *ICLR*, 2016.
- Oriol Vinyals, Meire Fortunato, and Navdeep Jaitly. Pointer networks. In Advances in Neural Information Processing Systems 28: Annual Conference on Neural Information Processing Systems 2015, December 7-12, 2015, Montreal, Quebec, Canada, pp. 2692–2700, 2015. URL http://papers.nips.cc/paper/5866-pointer-networks.
- Wojciech Zaremba, Tomas Mikolov, Armand Joulin, and Rob Fergus. Learning simple algorithms from examples. In *Proceedings of the 33nd International Conference on Machine Learning, ICML 2016, New York City, NY, USA, June 19-24, 2016*, pp. 421–429, 2016. URL http://jmlr.org/proceedings/papers/v48/zaremba16.html.

## A APPENDIX

#### A.1 PROGRAM SET FOR NON-RECURSIVE TOPOLOGICAL SORT

Program	Descriptions	Calls	Arguments
TOPOSORT	Perform topological	TRAVERSE,	NONE
	sort on graph	NEXT_START,	
		WRITE, MOVE	
TRAVERSE	Traverse graph until	CHECK_CHILD, EX-	NONE
	stack is empty	PLORE	
CHECK_CHILD	Check if a white	MOVE	NONE
	child exists; if so, set		
	$childList[v_{active}]$ to		
	point to it		
EXPLORE	Repeatedly traverse	STACK,	NONE
	subgraphs until stack	CHECK_CHILD,	
	is empty	WRITE, MOVE	
STACK	Interact with stack,	WRITE, MOVE	PUSH, POP
	either pushing or		
	popping		
NEXT_START	Move $p_{start}$ until	MOVE	NONE
	reaching a white		
	vertex. If a white		
	vertex is found, set		
	$p_{start}$ to point to it;		
	this signifies the start		
	of a traversal of a		
	new connected com-		
	ponent. If no white		
	vertex is found, the		
	entire execution is		
	terminated		
WRITE	Write a value either	NONE	Described below
	to environment (e.g.,		
	to color a vertex)		
	or variable (e.g., to		
	change the value of		
MOVE	$v_{active})$	NONE	7 11 1
MOVE	Move a pointer	NONE	Described below
	(e.g., $p_{start}$ or		
	$childList[v_{active}])$		
	up or down		

## **Argument Sets for WRITE and MOVE.**

**WRITE.** The WRITE operation has the following arguments:

*ARG\_1* (Main Action): COLOR\_CURR, COLOR\_NEXT, ACTIVE\_START, ACTIVE\_NEIGHB, ACTIVE\_STACK, SAVE, STACK\_PUSH, STACK\_POP, RESULT

COLOR\_CURR colors  $v_{active}$ , COLOR\_NEXT colors  $V_{active}$   $DAG[v_{active}][childList[v_{active}]]$ , ACTIVE\_START writes  $p_{start}$  to  $v_{active}$ , ACTIVE\_NEIGHB writes  $DAG[v_{active}][childList[v_{active}]]$  to  $v_{active}$ , ACTIVE\_STACK writes  $Q_{stack}(p_{stack})$  to  $v_{active}$ , SAVE writes  $v_{active}$  to  $v_{save}$ ,  $STACK\_PUSH$  pushes  $v_{active}$  to the top of the stack,  $STACK\_POP$  writes a null value to the top of the stack, and RESULT writes  $v_{active}$  to  $Q_{result}(p_{result})$ .

ARG\_2 (Auxiliary Variable): COLOR\_GREY, COLOR\_BLACK

COLOR\_GREY and COLOR\_BLACK color the given vertex grey and black, respectively.

## **MOVE.** The MOVE operation has the following arguments:

 $ARG\_1$  (Pointer):  $p_{result}$ ,  $p_{stack}$ ,  $p_{start}$ ,  $childList[v_{active}]$ ,  $childList[v_{save}]$ 

Note that the argument is the identity of the pointer, not what the pointer points to; in other words,  $ARG_{-}I$  can only take one of 5 values.

ARG\_2 (Increment or Decrement): UP, DOWN

#### A.2 TRACE-GENERATING FUNCTIONS FOR TOPOLOGICAL SORT

#### A.2.1 Non-Recursive Trace-Generating Functions

```
// Top level topological sort call
     TOPOSORT() {
       while (Q_{color}(p_{start}) is a valid color): // color invalid when all vertices explored
          WRITE (ACTIVE_START)
          WRITE (COLOR_CURR, COLOR_GREY)
          TRAVERSE()
         MOVE (p_{start}, 	ext{UP})
NEXT_START()
 8
10
11
12
     TRAVERSE() {
       CHECK CHILD()
13
14
15
       EXPLORE()
     CHECK_CHILD() {
       17
18
19
20
21
     EXPLORE() {
22
          if (Q_{color}(DAG[v_{active}][childList[v_{active}]]) is white): WRITE(COLOR_NEXT, COLOR_GREY)
23
24
25
26
27
            STACK (PUSH)
            WRITE (SAVE)
            WRITE (ACTIVE_NEIGHB)
28
29
            \texttt{MOVE}\,(childList[v_{save}], \texttt{ UP})
          else:
30
31
32
            WRITE (COLOR_CURR, COLOR_BLACK)
            WRITE (RESULT)
            \begin{array}{l} \text{MOVE}\left(p_{result}, \text{ UP}\right) \\ \text{if}\left(p_{stack} == 1\right): \end{array}
33
34
35
36
37
38
              break
            else:
              STACK (POP)
          CHECK_CHILD()
       while (true)
39
40
41
     STACK(op) {
       if (op == PUSH):
42
43
          WRITE (STACK_PUSH)
         MOVE (p_{stack}, \text{ UP})
45
46
       if (op == POP):
47
          WRITE (ACTIVE_STACK)
48
49
          WRITE (STACK_POP)
         MOVE (p_{stack}, DOWN)
50
51
52
53
54
       while (Q_{color}(p_{start})) is not white and is not invalid): // color invalid when all vertices explored
         MOVE (p_{start}, 	ext{UP})
```

## A.2.2 RECURSIVE TRACE-GENERATING FUNCTIONS

#### Altered Recursive Functions

```
// Top level topological sort call
      TOPOSORT() {
         if (Q_{color}(p_{start}) is a valid color): // color invalid when all vertices explored WRITE(ACTIVE_START)
            WRITE (COLOR_CURR, COLOR_GREY)
            TRAVERSE()
            \texttt{MOVE}\left(p_{start}, \texttt{UP}\right)
            NEXT_START()
            TOPOSORT() // Recursive Call
10
11
12
      CHECK_CHILD() {
         if (Q_{color}(DAG[v_{active}][childList[v_{active}]]) is not white and is not invalid): // color invalid when all children explored MOVE (childList[v_{active}], \text{ UP}) CHECK_CHILD() // Recursive Call
13
14
15
16
17
      EXPLORE() {
         if (Q_{color}(DAG[v_{active}][childList[v_{active}]]) is white): WRITE(COLOR_NEXT, COLOR_GREY)
19
20
21
22
            STACK (PUSH)
            WRITE (SAVE)
23
            WRITE (ACTIVE_NEIGHB)
24
25
            MOVE (childList[v_{save}], \text{ UP})
         else:
26
27
28
            WRITE (COLOR_CURR, COLOR_BLACK)
            WRITE (RESULT)
            \begin{aligned} & \text{MOVE}\left(p_{result}, \text{ UP}\right) \\ & \text{if}\left(p_{stack} == 1\right): \\ & \text{return} \end{aligned}
29
30
31
            else:
32
               STACK (POP)
         CHECK_CHILD()
EXPLORE() // Recursive Call
33
34
35
36
37
      NEXT_START() {
38
         if (Q_{color}(p_{start}) is not white and is not invalid): // color invalid when all vertices explored
39
            40
41
```

#### A.3 TRAINING SET FOR TOPOLOGICAL SORT

The single training problem input was a directed acyclic graph with 5 vertices corresponding to the following edge set: [(1, 2), (1, 5), (2, 4), (2, 5), (3, 5)].

## A.4 VERIFICATION SET FOR TOPOLOGICAL SORT

In order to verify the learned program is correct, it is sufficient to test on a graph of size 6 corresponding to the following edge set: [(1, 2), (1, 5), (2, 4), (2, 5), (3, 5)]. Note that there is a lone vertex (namely Vertex 6), not connected to any other vertices.

## A.5 Non-Recursive Quicksort

#### **Algorithm 4** Iterative Quicksort

```
1: Initialize an array A to sort and two empty stacks S_{lo} and S_{hi}.
 2: Initialize lo and hi to be 1 and n, where n is the length of A.
 3:
 4: function QUICKSORTWRAP(A, lo, hi)
 5:
        Push lo and hi onto S_{lo} and S_{hi}.
 6:
        QUICKSORT(A, lo, hi)
 8: function PARTITION(A, lo, hi)
 9:
        pivot = lo
10:
        for j \in [lo, hi-1]: do
            if A[j] \leq A[hi] then
11:
                swap A[pivot] with A[j]
12:
13:
                pivot = pivot + 1
14:
        swap A[pivot] with A[hi]
15:
        return pivot
16:
17: function QUICKSORT(A, lo, hi)
        while S_{lo} and S_{hi} are not empty: do
18:
19:
            Pop states off S_{lo} and S_{hi}, writing them to lo and hi.
20:
            p = PARTITION(A, lo, hi)
21:
            Push p + 1 and hi to S_{lo} and S_{hi}.
            Push lo and p-1 to S_{lo} and S_{hi}.
22:
```

## A.6 PROGRAM SET FOR QUICKSORT

Program	Descriptions	Calls	Arguments
QUICKSORT_WRAP	Wrapper function for	QUICKSORT, STACK	NONE
	non-recursive quick-		
	sort. Pushes ini-		
	tial $lo/hi$ onto stacks		
	and then launches the		
	main quicksort func-		
OT HOLICO DE	tion		× 11 1.1
QUICKSORT	Run the quicksort	Non-Recursive:	Implicitly: array
	routine in place for the array $A$ , for	PARTITION, RE- SET_PIVOT, RE-	A to sort, $lo$ , $hi$
	the array $A$ , for indices from $lo$ to $hi$	SET_J, STACK,	
	maices from to to no	WRITE	
		Recursive: same as	
		non-recusive version,	
		along with QUICK-	
		SORT	
PARTITION	Runs the partition	COMPSWAP_LOOP,	NONE
	function. At end,	MOVE_PIVOT_LO,	
	pointer $p_{pivot}$ is	MOVE_J_LO, SWAP	
	moved to the pivot		
COMPSWAP_LOOP	Runs the FOR loop	COMPSWAP, MOVE	NONE
	inside the partition function		
COMPSWAP	Compares	SWAP, MOVE	NONE
COMESWAF	$A[pivot] \leq A[j];$ if	SWAF, MOVE	NONE
	so, perform a swap		
	and increment $p_{pivot}$		
RESET_PIVOT	Resets $p_{pivot}$ to be-	MOVE	NONE
	ginning of array		
RESET_J	Resets $p_j$ to begin-	MOVE	NONE
	ning of array		
MOVE_PIVOT_LO	Moves $p_{pivot}$ to $lo$ in-	MOVE	NONE
	dex		
MOVEJLO	Moves $p_j$ to $lo$ index	MOVE	NONE
STACK	Pushes lo/hi states	WRITE, MOVE	Described below
	onto stacks $S_{lo}$ and		
	$S_{hi}$ according to		
	argument (described below)		
MOVE	Moves pointer one	NONE	Described below
111011	unit up or down	TIOTIL	Described below
SWAP	Swaps elements at	NONE	Described below
	given array indices		
WRITE	Write a value	NONE	Described below
	either to stack		
	(e.g., $Q_{stackLo}$ or		
	$Q_{stackHi}$ ) or to		
	pointer (e.g., to		
	change the value of		
	$p_{hi}$		

## Argument Sets for STACK, WRITE, MOVE, and SWAP.

**STACK.** The STACK operation has the following arguments:

*ARG\_1* (Operation): STACK\_PUSH\_ORI, STACK\_PUSH\_CALL1, STACK\_PUSH\_CALL2, STACK\_POP

STACK\_PUSH\_ORI pushes lo and hi to  $Q_{stackLo}$  and  $Q_{stackHi}$ . STACK\_PUSH\_CALL1 pushes lo and pivot-1 to  $Q_{stackLo}$  and  $Q_{stackHi}$ . STACK\_PUSH\_CALL2 pushes pivot+1 and hi to  $Q_{stackLo}$  and  $Q_{stackHi}$ . STACK\_POP pushes null values to  $Q_{stackLo}$  and  $Q_{stackHi}$ .

#### **WRITE.** The WRITE operation has the following arguments:

 $ARG_{-}1$  (Object to Write): ENV\_STACK\_LO, ENV\_STACK\_HI,  $p_{hi}$ ,  $p_{lo}$ 

ENV\_STACK\_LO and ENV\_STACK\_HI represent  $Q_{stackLo}(p_{stackLo})$  and  $Q_{stackHi}(p_{stackHi})$ , respectively.

 $ARG_2$  (Object to Copy): ENV\_STACK\_LO\_PEEK, ENV\_STACK\_HI\_PEEK,  $p_{hi}$ ,  $p_{lo}$ ,  $p_{pivot}-1$ ,  $p_{pivot}+1$ , RESET

ENV\_STACK\_LO\_PEEK and ENV\_STACK\_HI\_PEEK represent  $Q_{stackLo}(p_{stackLo}-1)$  and  $Q_{stackHi}(p_{stackHi}-1)$ , respectively. RESET represents a null value.

Note that the argument is the identity of the pointer, not what the pointer points to; in other words,  $ARG_{-}1$  can only take one of 4 values, and  $ARG_{-}2$  can only take one of 7 values.

#### **MOVE.** The MOVE operation has the following arguments:

```
ARG_{-}1 (Pointer): p_{stackLo}, p_{stackHi}, p_j, p_{pivot}
```

Note that the argument is the identity of the pointer, not what the pointer points to; in other words,  $ARG_{-}I$  can only take one of 4 values.

ARG\_2 (Increment or Decrement): UP, DOWN

#### **SWAP.** The SWAP operation has the following arguments:

```
ARG_{-}I (Swap Object 1): p_{pivot}

ARG_{-}2 (Swap Object 2): p_{hi}, p_{j},
```

#### A.7 TRACE-GENERATING FUNCTIONS FOR QUICKSORT

## A.7.1 Non-Recursive Trace-Generating Functions

```
Initialize p_{lo} to 1 and p_{hi} to n (length of array) // Wrapper function for quicksort
      QUICKSORT_WRAP() {
   STACK(STACK_PUSH_ORI)
          QUICKSORT()
 6
7
       QUICKSORT() {
          \begin{aligned} & \text{while } & (p_{stackLo} \neq 1): \\ & \text{if } & (Q_{stackLo}(p_{stackLo} - 1) < Q_{stackHi}(p_{stackHi} - 1)): \\ & \text{STACK} & (\text{STACK\_POP}) \end{aligned}
11
12
              else:
13
14
15
                 WRITE (p_{hi}, ENV_STACK_HI_PEEK)
                 WRITE (p_{lo}, \text{ ENV\_STACK\_LO\_PEEK})
STACK (STACK_POP)
16
17
18
                 PARTITION()
                 STACK (STACK_PUSH_CALL2)
                 STACK (STACK_PUSH_CALL1)
19
                 RESET_PIVOT()
20
                 RESET_J()
21
22
23
      PARTITION() {
24
          MOVE_PIVOT_LO()
          MOVE_J_LO()
COMPSWAP_LOOP()
25
26
27
          \texttt{SWAP}\left(p_{pivot}, p_{hi}\right)
28
29
30
     COMPSWAP LOOP() {
31
          while (p_j \neq p_{hi}):
              COMPSWAP()
32
33
              MOVE (p_j , UP)
34
```

```
36
37
       COMPSWAP() {
           if (A[p_j] \leq A[p_{hi}]):
38
               \texttt{SWAP}\,(p_{pivot},p_{j})
39
               MOVE (p_{pivot}, \ \text{UP})
40
       }
41
42
       RESET_PIVOT() {
           while (p_{pivot} \neq 1):
43
44
              MOVE (p_{pivot}, \text{ DOWN})
45
46
47
       RESET_J() {
48
           while (p_j \neq 1):
49
               MOVE (p_j , DOWN)
50
51
       MOVE_PIVOT_LO() {
53
            while (p_{pivot} \neq p_{lo}):
54
               MOVE (p_{pivot}, \ \text{UP})
55
56
57
       MOVE_J_LO() {
58
           while (p_j \neq p_{lo}):
59
               MOVE (p_j , UP)
60
61
62
       STACK(op) {
          if (op == STACK_PUSH_ORI): WRITE (ENV_STACK_LO, p_{lo}) WRITE (ENV_STACK_HI, p_{hi})
63
64
65
               \begin{array}{l} \text{MOVE}\left(p_{stackLo}, \text{ UP}\right) \\ \text{MOVE}\left(p_{stackHi}, \text{ UP}\right) \end{array}
66
67
68
69
           if (op == STACK_PUSH_CALL1):
               WRITE (ENV_STACK_HI, p_{lo})
WRITE (ENV_STACK_HI, p_{pivot}-1)
70
71
72
73
74
               \texttt{MOVE}\left(p_{stackLo}, \texttt{UP}\right)
               \texttt{MOVE}\left(p_{stackHi}, \text{ UP}\right)
75
76
            if (op == STACK_PUSH_CALL2):
               WRITE (ENV_STACK_LO, p_{pivot}+1)
               WRITE(ENV_STACK_HI, p_{hi})
77
78
               \begin{array}{ll} \text{MOVE} \left( p_{stackLo}, \text{ UP} \right) \\ \text{MOVE} \left( p_{stackHi}, \text{ UP} \right) \end{array}
79
80
            if (op == STACK_POP):
81
82
               WRITE (ENV_STACK_LO, RESET)
               WRITE (ENV_STACK_HI, RESET) MOVE (p_{stackLo}, DOWN)
83
85
               \texttt{MOVE}\left(p_{stackHi}, \; \texttt{DOWN}\right)
86
```

#### A.7.2 RECURSIVE TRACE-GENERATING FUNCTIONS

### Altered Recursive Functions

```
No QUICKSORT_WRAP
         Initialize p_{lo} to 1 and p_{hi} to n (length of array)
 4
            \begin{aligned} &\text{if } &(Q_{stackLo}(p_{stackLo}-1) < Q_{stackHi}(p_{stackHi}-1)): \\ &\text{partition()} \\ &\text{STACK}(&\text{STACK}\_\text{PUSH\_CALL2}) \end{aligned}
                 STACK (STACK_PUSH_CALL1)
                 RESET_PIVOT()
10
11
                 RESET_J()
                 \begin{array}{l} \texttt{WRITE}\left(p_{hi}\text{, ENV\_STACK\_HI\_PEEK}\right) \\ \texttt{WRITE}\left(p_{lo}\text{, ENV\_STACK\_LO\_PEEK}\right) \end{array}
12
13
14
                 {\tt QUICKSORT}\,()\ //\ \textbf{Recursive Call}
                 STACK (STACK_POP)
                 \begin{array}{ll} \text{WRITE} \left(p_{hi}, \text{ ENV\_STACK\_HI\_PEEK}\right) \\ \text{WRITE} \left(p_{lo}, \text{ ENV\_STACK\_LO\_PEEK}\right) \\ \text{QUICKSORT}() \ // \ \textbf{Recursive Call} \end{array}
15
16
17
18
                 STACK (STACK_POP)
19
20
21
22
        COMPSWAP_LOOP() {
            if (p_j \neq p_{hi}):
                 COMPSWAP()
23
24
                 MOVE (p_j , UP)
25
                 COMPSWAP_LOOP() // Recursive Call
26
27
28
        RESET_PIVOT() {
29
            if (p_{pivot} \neq 1):
30
                MOVE (p_{pivot}, \text{ DOWN})
```

```
31
32
33
                 RESET_PIVOT() // Recursive Call
34
35
       RESET_J() { if (p_j \neq 1):
36
                MOVE (p_j , DOWN)
37
38
39
40
41
                RESET_J() // Recursive Call
        MOVE_PIVOT_LO() {
            if (p_{pivot} \neq p_{lo}):
42
                MOVE (p_{pivot}, \text{UP})

MOVE_PIVOT_LO() // Recursive Call
43
44
45
46
47
       \begin{array}{l} \mbox{MOVE\_J\_LO\,() \ \{} \\ \mbox{if } (p_j \neq p_{lo}): \\ \mbox{MOVE}(p_j, \mbox{ UP}) \\ \mbox{MOVE\_J\_LO\,() // Recursive Call} \end{array}
48
49
50
```

## A.8 VERIFICATION SET FOR QUICKSORT

In order to verify the learned program is correct, it is sufficient to test on the following set of 3 arrays:  $\{[2,8,9,4,5,9,2,2,4],[3,4,9,7,4,4,9,1,5],[1,2,3,4,5,6,7]\}.$